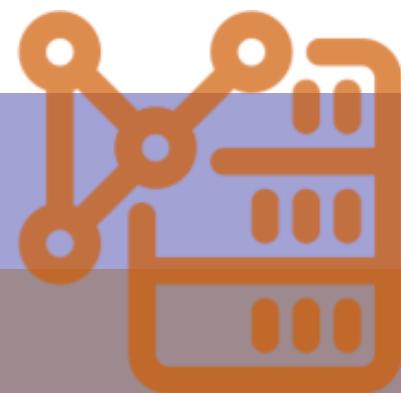




سہیل عمران



Dr. SOHAIL IMRAN



Intro

The concept of NOSQL databases gained popularity among Internet giants like Google, Yahoo, Facebook, and Amazon, which deal with huge volumes of data.
Carl Strozz introduced the NOSQL concept in 1998.

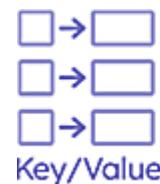
NOSQL stands for "**Not Only Structured Query Language**".

NOSQL should not be misleading:
the approach does not prohibit Structured Query Language (SQL)

It is a non-relational DMS, that does not require a fixed schema, avoids joins, and is easy to scale.

The purpose of using a NOSQL database is for distributed data stores with humongous data storage needs. NOSQL is used for Big data and real-time web apps.

Traditional **RDBMS** uses **SQL** syntax to store and retrieve data for further insights. Instead, a NOSQL database system encompasses a wide range of database technologies that can store structured, semi-structured, unstructured and polymorphic data.



Key/Value



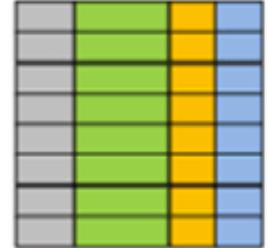
Graph



Column



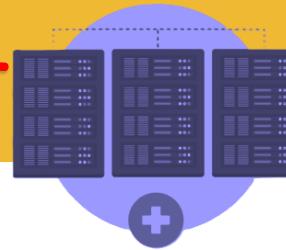
Document



Relational



scale-up Vs scale-out



Problem with RDMS:

The system response time becomes slow when you use RDBMS for massive volumes of data.

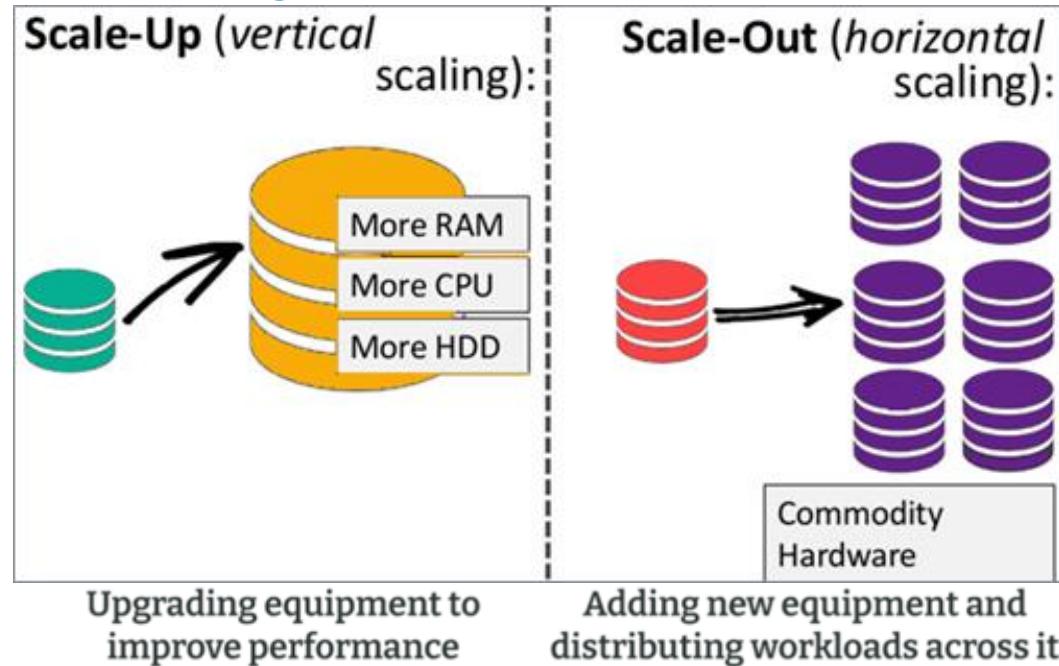
Solution:

To resolve this problem, we could "scale up" our systems by upgrading our existing hardware.

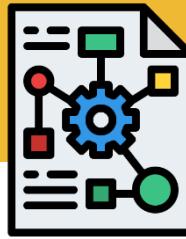
This process is **expensive**.

The alternative to this issue is to distribute the database load across multiple hosts whenever the load increases. This method is known as "**scaling out**."

A NOSQL database is non-relational, so it scales out better than relational databases, as they are designed with web applications in mind.



NoSQL



features

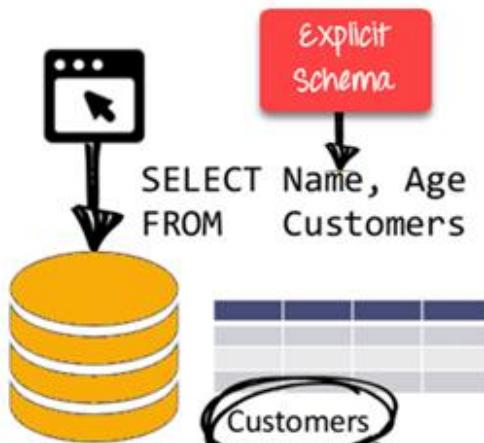
Non-relational

- NOSQL databases never follow the relational model
- Never provide tables with flat fixed-column records
- Work with self-contained aggregates or BLOBs
- Doesn't require object-relational mapping and data normalization
- No complex features like query languages, query planners, referential integrity joins, ACID

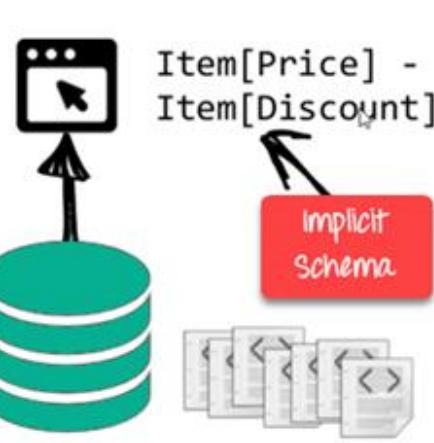
Schema-free

- NOSQL databases are either schema-free or have relaxed schemas
- Do not require any sort of definition of the schema of the data
- Offers heterogeneous structures of data in the same domain

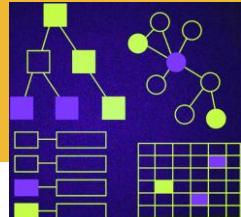
RDBMS:



NoSQL DB:



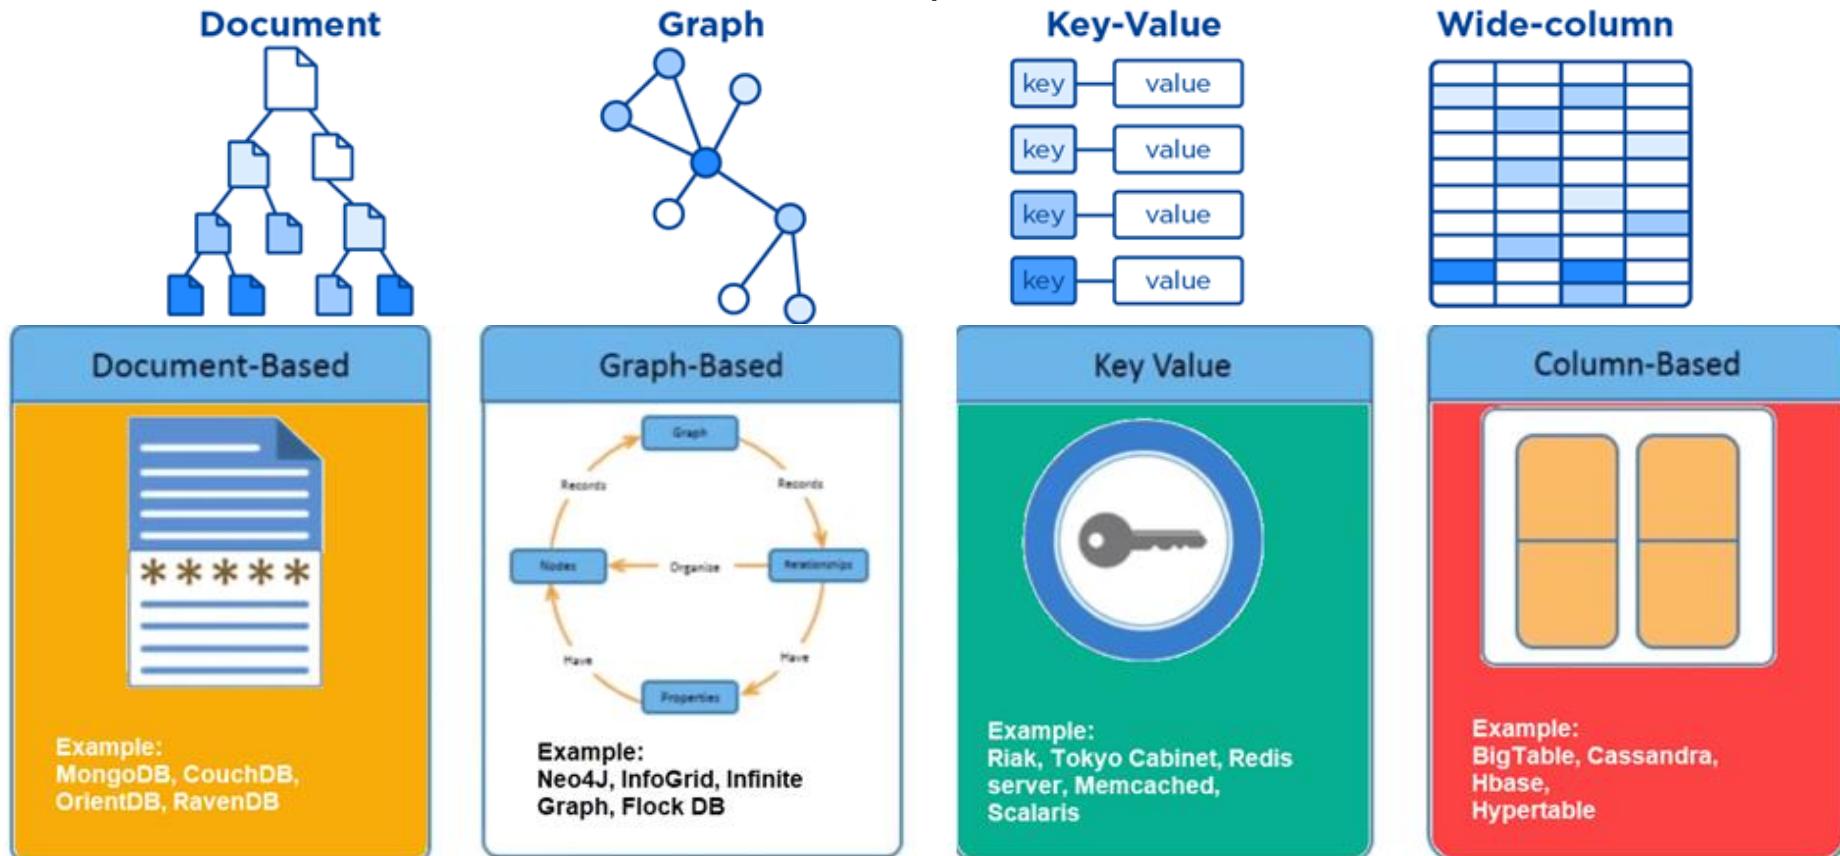
An **implicit schema** means that a document can be stored as it is without having to define a **schema** for it and without checking that it conforms to a **schema**.

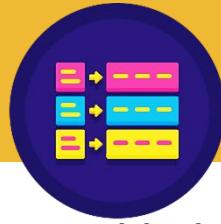


types



There are mainly four categories of NOSQL databases.
Each of these categories has its unique attributes and limitations.
No specific database is better to solve all problems.
One should select a database based on one's product needs.





key value pair-based DB



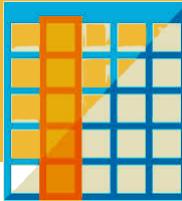
Data is stored in key/value pairs. It is designed in such a way to handle lots of data and a heavy load.

Key-value pair storage databases store data as a hash table where each key is unique, and the value can be a JSON, BLOB(Binary Large Objects), string, etc.

It is one of the most basic types of NOSQL databases. This kind of NOSQL database is used as a collection, dictionaries, associative arrays, etc. Key value stores help the developer to store schema-less data. They work best for shopping cart contents.

Redis, Dynamo, Riak are some examples of key-value store Databases. They are all based on Amazon's Dynamo paper.

Key	Value
Name	Joe Bloggs
Age	42
Occupation	Stunt Double
Height	175cm
Weight	77kg



column-oriented database



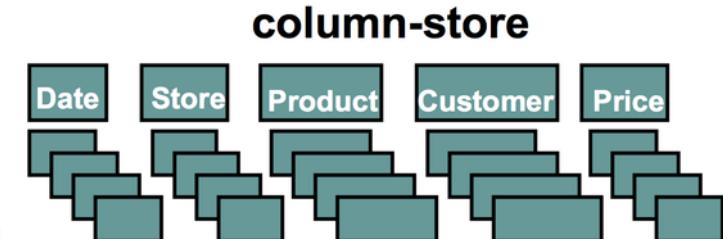
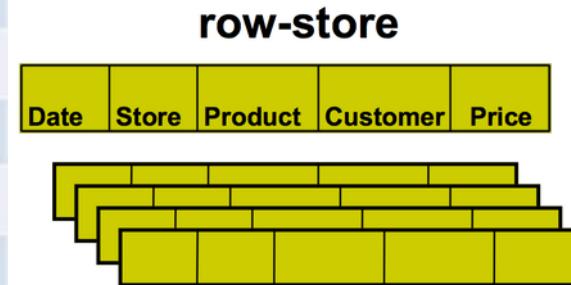
Column-oriented databases work on columns and are based on BigTable paper by Google. Every column is treated separately. Values of single column databases are stored contiguously.

They deliver high performance on aggregation queries like SUM, COUNT, AVG, MIN etc. as the data is readily available in a column.

Column-based NOSQL databases are widely used to manage data warehouses, business intelligence, CRM, Library card catalogs,

HBase, Cassandra, HBase, Hypertable are examples of column based database.

ColumnFamily			
Row Key	Column Name		
	Key	Key	Key
Value	Value	Value	Value
	Column Name		
Key	Key	Key	Key
	Value	Value	Value





document-oriented DB store

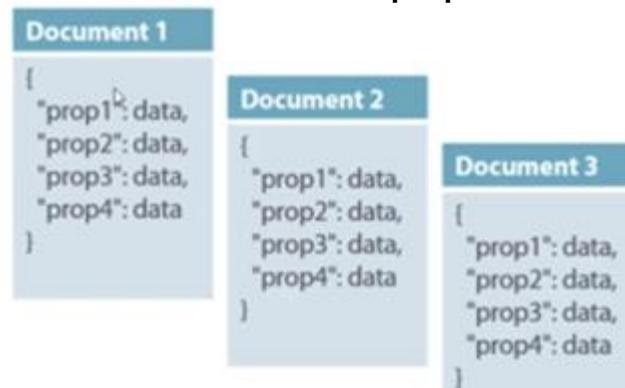


Document-Oriented NOSQL DB stores and retrieves data as a key value pair but the value part is stored as a document. The document is stored in JSON or XML formats. The value is understood by the DB and can be queried.

The document type is mostly used for CMS systems, blogging platforms, real-time analytics & e-commerce applications. It should not use for complex transactions which require multiple operations or queries against varying aggregate structures.

Amazon SimpleDB, CouchDB, MongoDB, Riak, Lotus Notes, are popular Document originated DBMS systems.

Col1	Col2	Col3	Col4
Data	Data	Data	Data
Data	Data	Data	Data
Data	Data	Data	Data



In this diagram, you can see we have rows and columns, and in the right, we have a document database which has a similar structure to JSON. Now for the relational database, you have to know what columns you have and so on. However, for a document database, you have data store like JSON object. You do not require to define which make it flexible.



graph-based database

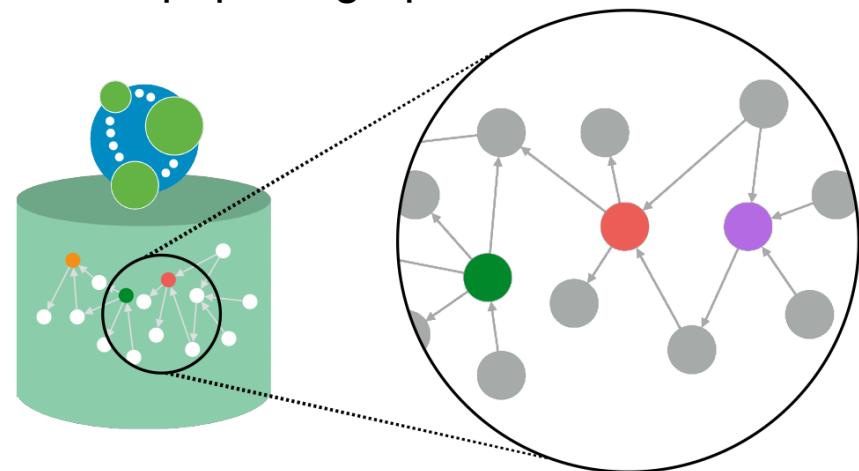
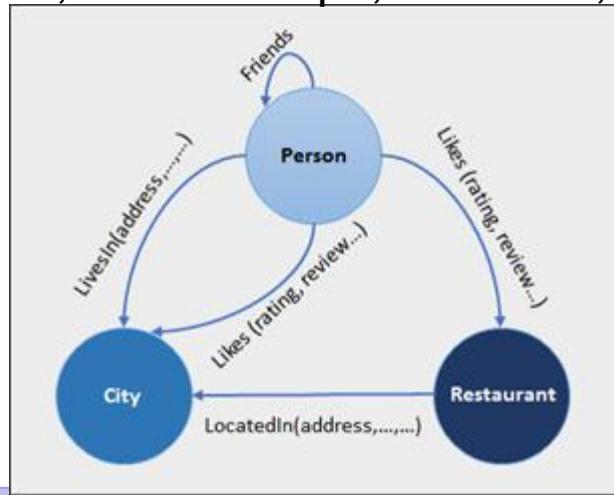


A graph type database stores entities as well the relations amongst those entities. The entity is stored as a node with the relationship as edges. An edge gives a relationship between nodes. Every node and edge has a unique identifier.

Compared to a relational database where tables are loosely connected, a Graph database is a multi-relational in nature. Traversing relationship is fast as they are already captured into the DB, and there is no need to calculate them.

Graph base database mostly used for social networks, logistics, spatial data.

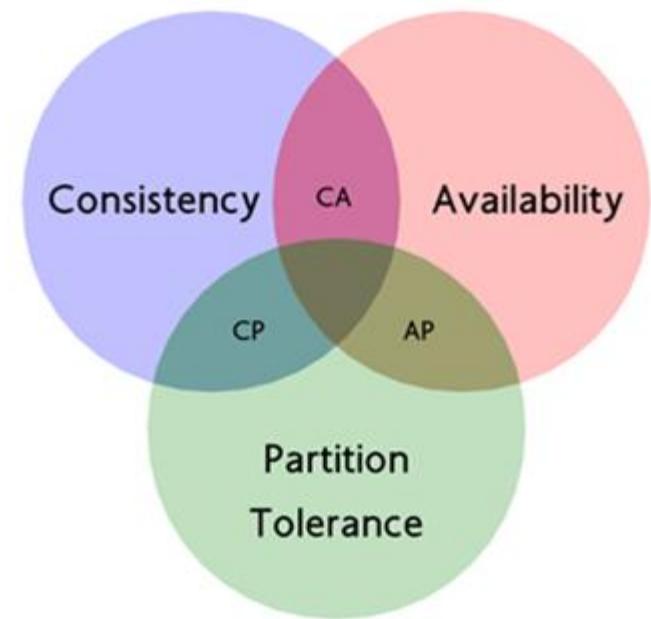
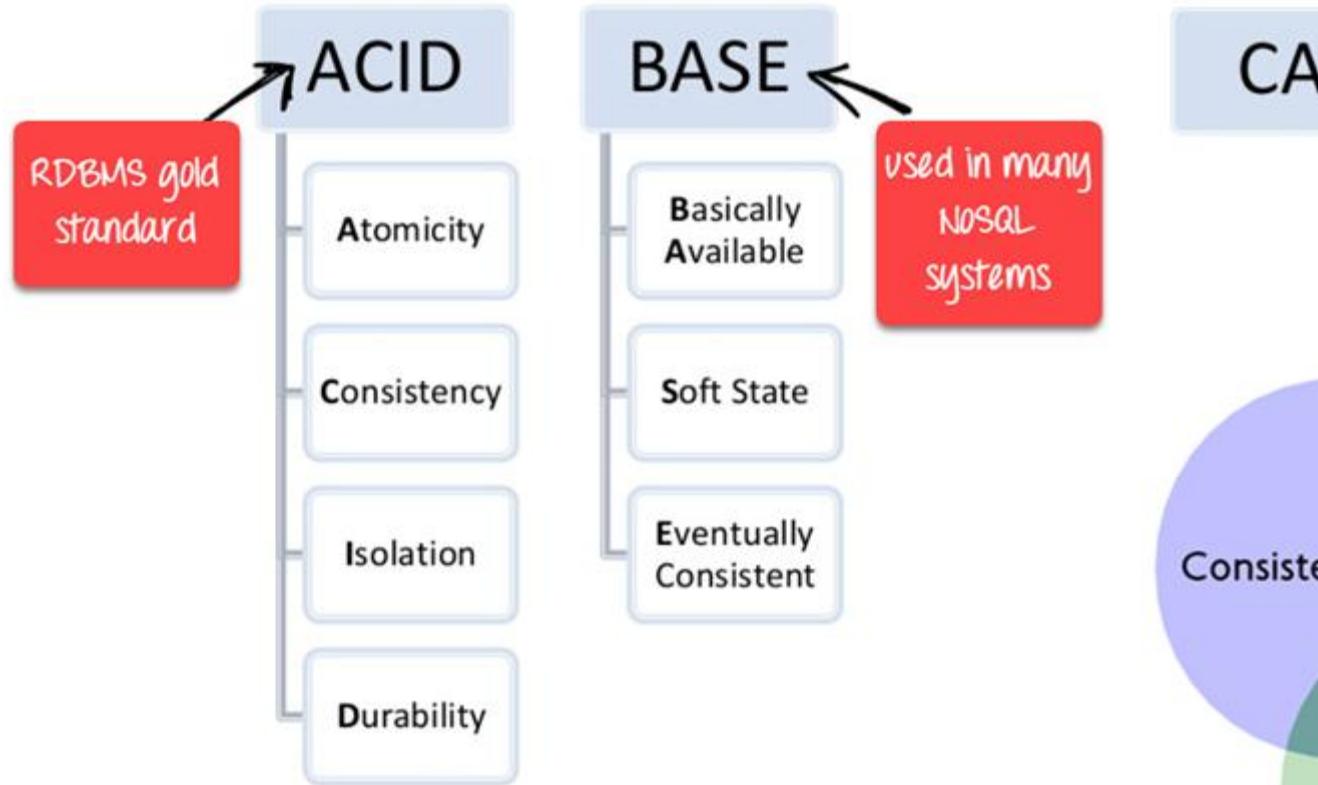
Neo4J, Infinite Graph, OrientDB, FlockDB are some popular graph-based databases.



NoSQL



A B C of data stores

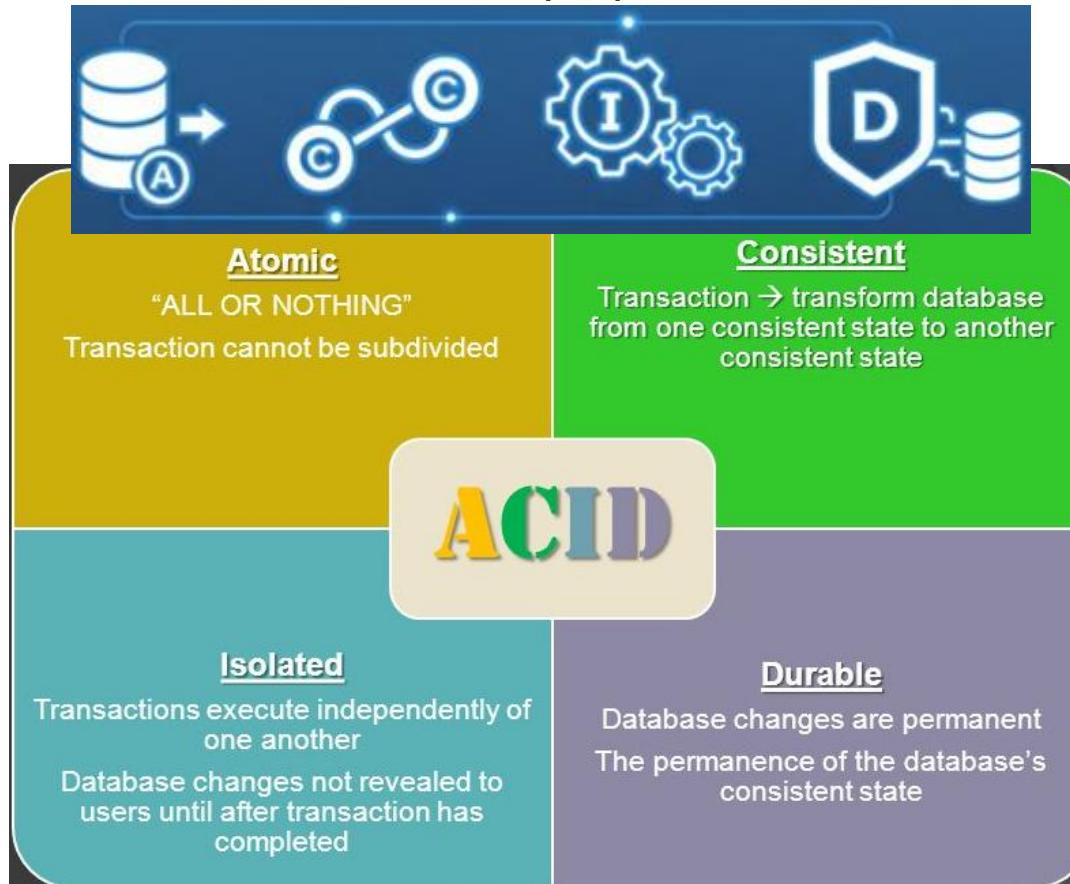




ACID properties



A **transaction** is a single logical unit of work that accesses and possibly modifies the contents of a database. Transactions access data using read and write operations. In order to maintain consistency in a database, before and after the transaction, certain properties are followed. These are called **ACID** properties.





BASE



BASE

= Basically Available, Soft state, Eventual consistency

If a node fails,
part of the data
will not be
available, but the
entire data layer
stays operational

The state of the
system may change
over time, even
without input

The system
becomes consistent
at some later time

Consistency model weaker than

ACID

Thus, changes made to any data item on one machine has to be propagated to other replicas. Data replication may not be instantaneous as some copies will be updated immediately while others in due course of time.

These copies may be mutually, but in due course of time, they become consistent. Hence, the name eventual consistency.

NoSQL



CAP theorem is also called brewer's theorem. It states that it is impossible for a distributed data store to offer more than two out of three guarantees.

Network Problem might stop the system
SQL/RDBMS

Consistency

All clients see the same view of data, even right after update or delete

Pick two

There is a risk of some data becoming unavailable
noSQL/column

Availability

All clients can find a replica of data, even in case of partial node failures

CA

Partitioning

The system continues to work as expected, even in presence of partial network failure

AP

Clients may read inconsistent data
noSQL/document or key/value



NoSQL



NoSQL Vs. SQL

NoSQL databases use various data models optimized for different use cases.	SQL databases use a relational data model based on tables with predefined columns and rows.
Uses query languages or APIs optimized for specific data models.	Uses a standardized language called Structured Query Language (SQL) to manipulate data.
NoSQL databases are horizontally scalable and can distribute data across multiple servers and nodes.	SQL databases are vertically scalable.
May not be fully ACID compliant	SQL databases are designed to be ACID compliant
Can provide faster performance for unstructured and non-relational data access	Can provide high performance for structured queries and data manipulation
NoSQL databases can be more cost-effective for specific use cases	SQL databases can be more expensive than NoSQL databases



NoSQL



Challenges of using Wrong Database

:neo4j



Graph queries can take many lines of SQL and slow to run



Running free text queries on SQL databases is often complicated and again can be slow



NoSQL databases often good at finding data based on key but cannot provide multi-field querying of an SQL database



In-memory DB fast so long as data can be stored in memory! What about large datasets?



Don't Choose One Database Choose Them All! - Version 1.0 | May 2017

Copyright © Capgemini 2017. All Rights Reserved

4

NoSQL



Why Not just use Multiple Databases?

:neo4j



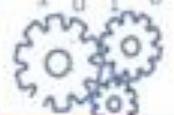
Data Science Technologies

- > 1000101101010110100010
- > 1000101111010110010010
- > 1011000101001101010101



Decouple Data and Analytics layers for analytics tool flexibility.

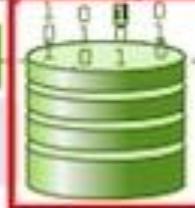
Database APIs



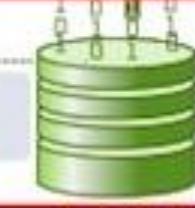
Most DS languages have good API support

Data Layer

Graph version



Lucene version



NoSQL version



Push as much data processing down into database.



Have several slave databases.

Capgemini
CONSULTING SERVICES

Don't Choose One Database Choose Them All! - Session 1.0 | May 2017

Copyright © Capgemini 2017. All Rights Reserved

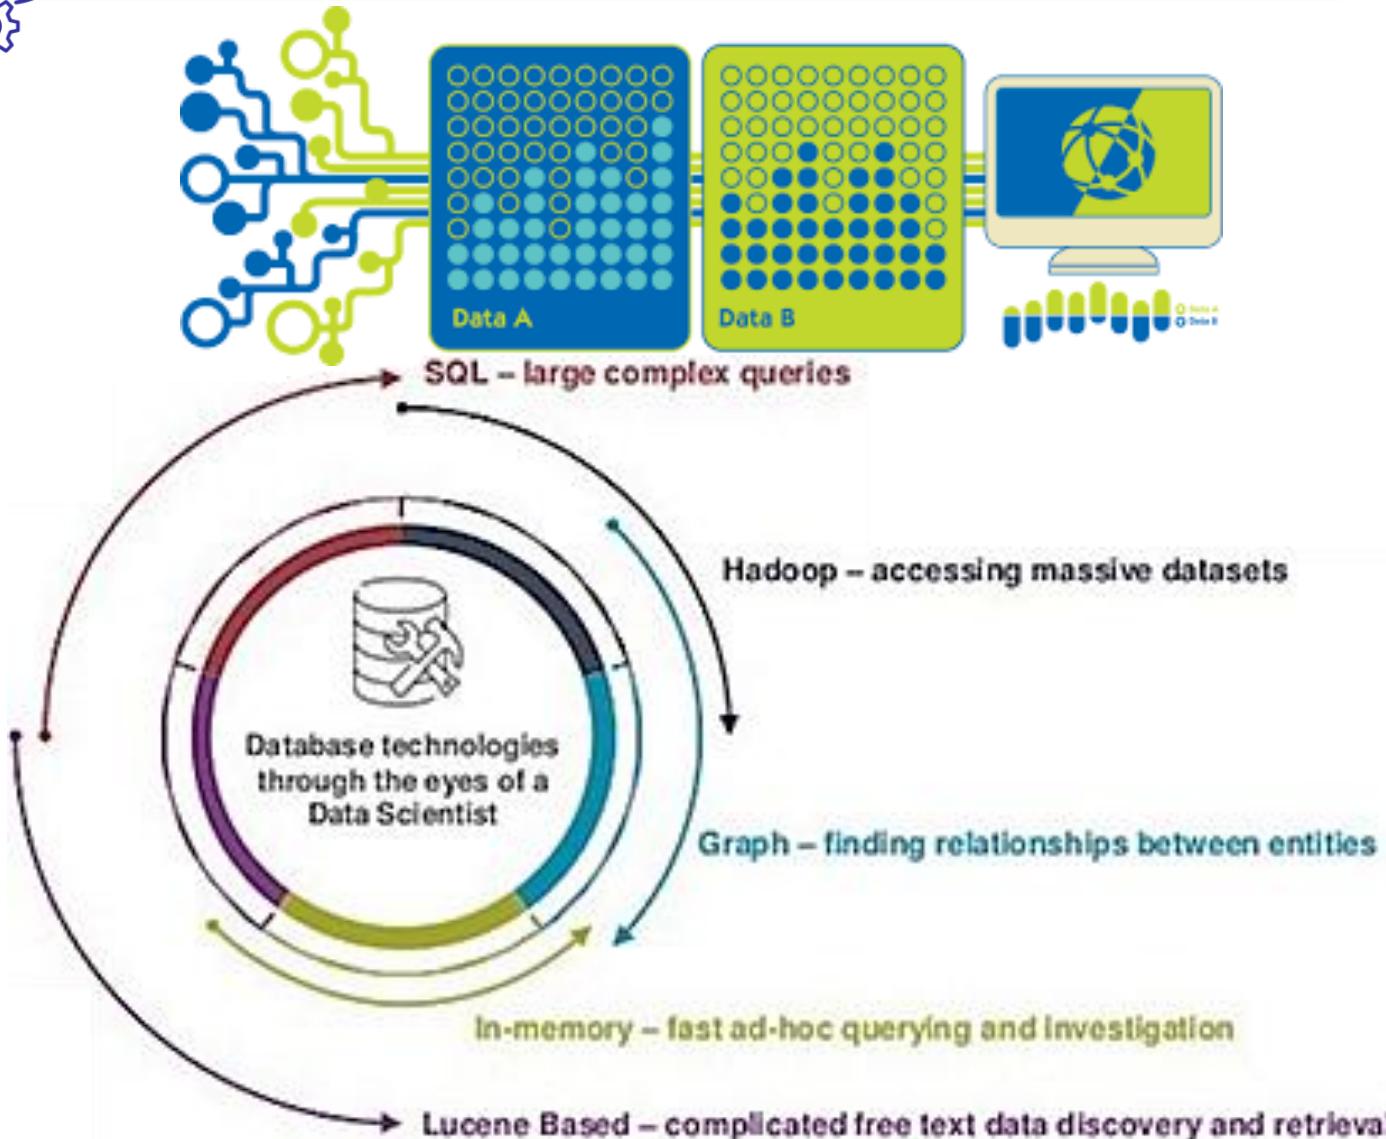
5

NoSQL



Big Database Solutions

:neo4j



NoSQL



Case Study: Car Insurance Fraud (CIF)



Credit Score	Bad Apps @ Address	Ave Annual Income
90	0	30,000
60	0	45,000
67	0	20,000
84	2	60,000
34	5	10,000

 Below the table is a bulleted list:

- Fast joins of multiple large tables
- Complex WHERE conditions on those joins

Applicant	Links To Bad Applicants
Jon	0
Jim	0
Joan	2
Janet	0
Jim Bob	1

- Graph queries less code & faster than same in SQL
- Out of the box graph queries

NoSQL

```

192.168.192.01 - [22/Dec/2015:21:10:20 -0400] "GET / HTTP/1.1" 200 639 "www.myata.com/app_page1"
"Mozilla/5.0 (compatible; MSE/6.0; Windows NT 6.1;)" */
192.168.192.01 - [22/Dec/2015:21:11:40 -0400] "GET /app_index.cgi HTTP/1.1" 200 867 "www.myata.com/app_page2"
"Mozilla/5.0 (Windows NT 6.1; rv:3.6) Gecko/20100101 Firefox/3.6" */
192.168.192.01 - [22/Dec/2002:21:12:10 -0400] "GET /app_index.cgi HTTP/1.1" 200 3900 "www.myata.com/app_page2"
"Mozilla/4.0 (compatible; MSE ...)" */

```

- Process large unstructured web logs
- Extract data and apply behaviour model

Applicant	Behaviour Normalcy
Jon	0.99
Jim	0.95
Joan	0.87
Janet	0.56
Jim Bob	0.82



Text

I am the best applicant ever, I promise, so no need to waste your time looking at my previous five convictions.

- Advanced text search
- Convert text into structured data

Applicant	Conviction	Fraud	Promise
Jon	0	0	0
Jim	0	0	1
Joan	0	1	0
Janet	1	1	1
Jim Bob	0	0	0

Guru Chaitanya Database Oracle
Consultant



Case Study: CIF - bringing it Together

neo4j



Applicant	Credit Score	Bad Apps @ Address	Ave Annual Income	Behaviour Normalcy	Conviction	Fraud	Promise
Jon	90	0	30,000	0.99	0	0	0
Jim	60	0	45,000	0.95	0	0	1
Joan	67	0	20,000	0.87	0	1	0
Janet	84	2	60,000	0.56	1	1	1
Jim Bob	34	5	10,000	0.82	0	0	0

Much richer base for Insight Generation

Value of data substantially increased by using different data base technologies.

Capgemini
Simplifying Complexity

Don't Choose One Database Choose Them All! - Version 7.0 | May 2017

Copyright © Capgemini 2017. All Rights Reserved

NoSQL