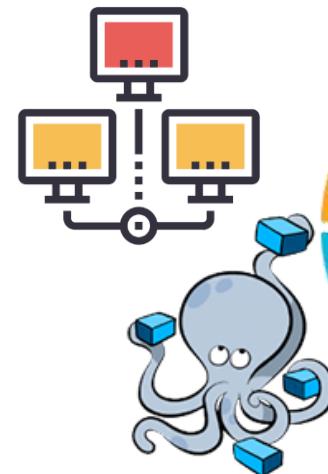


ORACLE®
Certified Professional



Microsoft Certified
Professional

سہیل عمران



BigData Analytics

Dr.SOHAİL IMRAN



Hadoop on Docker

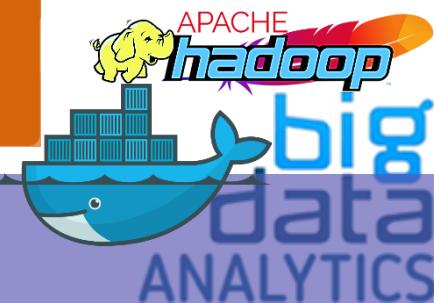


update



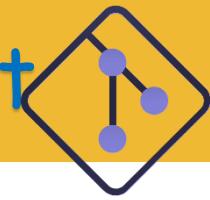
sudo apt update

```
sohailimran@DESKTOP-9VAOBA3:~$ sudo apt update
[sudo] password for sohailimran:
Hit:1 http://archive.ubuntu.com/ubuntu noble InRelease
Get:2 http://security.ubuntu.com/ubuntu noble-security InRelease [126 kB]
Get:3 http://archive.ubuntu.com/ubuntu noble-updates InRelease [126 kB]
Get:4 http://archive.ubuntu.com/ubuntu noble-backports InRelease [126 kB]
Get:5 http://archive.ubuntu.com/ubuntu noble-updates/main amd64 Packages [1585 kB]
Get:6 http://security.ubuntu.com/ubuntu noble-security/main amd64 Packages [1306 kB]
Get:7 http://archive.ubuntu.com/ubuntu noble-updates/main Translation-en [299 kB]
Get:8 http://archive.ubuntu.com/ubuntu noble-updates/main amd64 Components [175 kB]
Get:9 http://archive.ubuntu.com/ubuntu noble-updates/main amd64 c-n-f Metadata [15.7 kB]
Get:10 http://archive.ubuntu.com/ubuntu noble-updates/universe amd64 Packages [1498 kB]
Get:11 http://archive.ubuntu.com/ubuntu noble-updates/universe amd64 Components [378 kB]
Get:12 http://archive.ubuntu.com/ubuntu noble-updates/universe amd64 c-n-f Metadata [31.4 kB]
Get:13 http://archive.ubuntu.com/ubuntu noble-updates/restricted amd64 Packages [2244 kB]
Get:14 http://security.ubuntu.com/ubuntu noble-security/main Translation-en [215 kB]
Get:15 http://security.ubuntu.com/ubuntu noble-security/main amd64 Components [21.6 kB]
Get:16 http://security.ubuntu.com/ubuntu noble-security/main amd64 c-n-f Metadata [9336 B]
Get:17 http://security.ubuntu.com/ubuntu noble-security/universe amd64 Packages [906 kB]
Get:18 http://archive.ubuntu.com/ubuntu noble-updates/restricted Translation-en [509 kB]
Get:19 http://archive.ubuntu.com/ubuntu noble-updates/restricted amd64 Components [212 B]
Get:20 http://archive.ubuntu.com/ubuntu noble-updates/multiverse amd64 Components [940 B]
Get:21 http://archive.ubuntu.com/ubuntu noble-backports/main amd64 Components [7156 B]
Get:22 http://archive.ubuntu.com/ubuntu noble-backports/universe amd64 Components [11.0 kB]
Get:23 http://archive.ubuntu.com/ubuntu noble-backports/restricted amd64 Components [212 B]
Get:24 http://archive.ubuntu.com/ubuntu noble-backports/multiverse amd64 Components [212 B]
Get:25 http://security.ubuntu.com/ubuntu noble-security/universe amd64 Components [52.3 kB]
Get:26 http://security.ubuntu.com/ubuntu noble-security/universe amd64 c-n-f Metadata [19.4 kB]
Get:27 http://security.ubuntu.com/ubuntu noble-security/restricted amd64 Packages [2140 kB]
Get:28 http://security.ubuntu.com/ubuntu noble-security/restricted Translation-en [486 kB]
Get:29 http://security.ubuntu.com/ubuntu noble-security/restricted amd64 Components [208 B]
Get:30 http://security.ubuntu.com/ubuntu noble-security/multiverse amd64 Components [208 B]
Fetched 12.3 MB in 26s (478 kB/s)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
64 packages can be upgraded. Run 'apt list --upgradable' to see them.
sohailimran@DESKTOP-9VAOBA3:~$ |
```





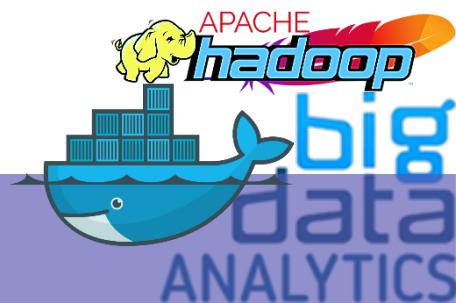
install git



```
curl --version  
sudo apt install git
```

docker desktop PERSONAL

```
sohailimran@DESKTOP-9VA0BA3:~$ curl --version  
curl 8.5.0 (x86_64-pc-linux-gnu) libcurl/8.5.0 OpenSSL/3.0.13 zlib/1.3 brotli/1.1.0 zstd/1.5.5 libidn2/2  
.3.7 libpsl/0.21.2 (+libidn2/2.3.7) libssh/0.10.6/openssl/zlib nghttp2/1.59.0 librtmp/2.3 OpenLDAP/2.6.7  
Release-Date: 2023-12-06, security patched: 8.5.0-2ubuntu10.6  
Protocols: dict file ftp ftps gopher gophers http https imap imaps ldap ldaps mqtt pop3 pop3s rtmp rtsp  
scp sftp smb smbs smtp smtps telnet tftp  
Features: alt-svc AsynchDNS brotli GSS-API HSTS HTTP2 HTTPS-proxy IDN IPv6 Kerberos Largefile libz NTLM  
PSL SPNEGO SSL threadsafe TLS-SRP UnixSockets zstd  
sohailimran@DESKTOP-9VA0BA3:~$ sudo apt install git  
[sudo] password for sohailimran:  
Reading package lists... Done  
Building dependency tree... Done  
Reading state information... Done  
git is already the newest version (1:2.43.0-1ubuntu7.3).  
git set to manually installed.  
0 upgraded, 0 newly installed, 0 to remove and 64 not upgraded.  
sohailimran@DESKTOP-9VA0BA3:~$ |
```





git clone



```
git clone https://github.com/big-data-europe/docker-hadoop
```

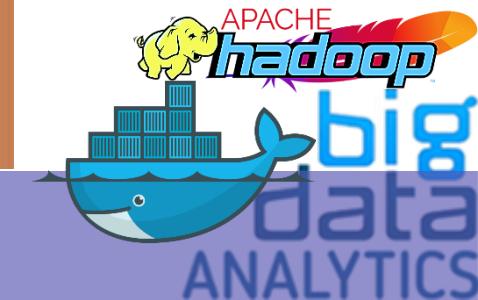
```
sohailimran@DESKTOP-9VAO ~
```

```
sohailimran@DESKTOP-9VAOBA3:~$ git clone https://github.com/big-data-europe/docker-hadoop
Cloning into 'docker-hadoop'...
remote: Enumerating objects: 539, done.
remote: Counting objects: 100% (189/189), done.
remote: Compressing objects: 100% (23/23), done.
remote: Total 539 (delta 169), reused 166 (delta 166), pack-reused 350 (from 1)
Receiving objects: 100% (539/539), 108.00 KiB | 654.00 KiB/s, done.
Resolving deltas: 100% (251/251), done.
sohailimran@DESKTOP-9VAOBA3:~$ ls
Dockerfile app2.py docker-hadoop
```

```
cd docker-hadoop
```

```
sohailimran@DESKTOP-9VAO ~
```

```
sohailimran@DESKTOP-9VAOBA3:~$ cd docker-hadoop
sohailimran@DESKTOP-9VAOBA3:~/docker-hadoop$ ls
Makefile      datanode          hadoop.env      nginx        submit
README.md    docker-compose-v3.yml  historyserver  nodemanager
base         docker-compose.yml   namenode       resourcemanager
```





docker-compose.yml



```
sohailimran@DESKTOP-9VAOBA3:~/docker-hadoop$ ls
Makefile    base        docker-compose-v3.yml  hadoop.env      namenode      nodemanager      submit
README.md   datanode    docker-compose.yml    historyserver  nginx        resourcemanager
```

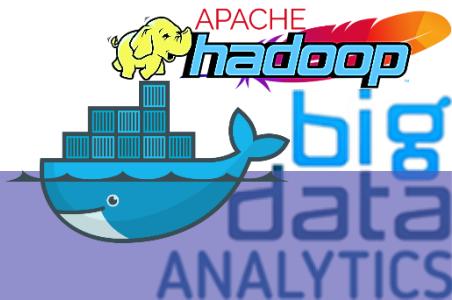
sudo nano docker-compose.yml

```
sohailimran@DESKTOP-9VAOBA3:~/docker-hadoop$ sudo nano docker-compose.yml
```

```
GNU nano 7.2                               docker-compose.yml *
version: "3"
services:
  namenode:
    image: bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8
    container_name: namenode
    restart: always
    ports:
      - 9870:9870
      - 9000:9000
      - 9864:9864
    volumes:
      - hadoop_namenode:/hadoop/dfs/name
    environment:
      - CLUSTER_NAME=docker-hadoop
    env_file:
      - ./hadoop.env
```

```
datanode:
  image: bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8
  container_name: datanode
  ports:
    - 9865:9864
  restart: always
  volumes:
    - hadoop_datanode:/hadoop/dfs/data
  environment:
    SERVICE_PRECONDITION: "namenode:9870"
  env_file:
    - ./hadoop.env
```

1. Ctrl + O
2. Enter
3. Ctrl + X





docker compose version



sohailimran@SI: ~/docker-hadoop



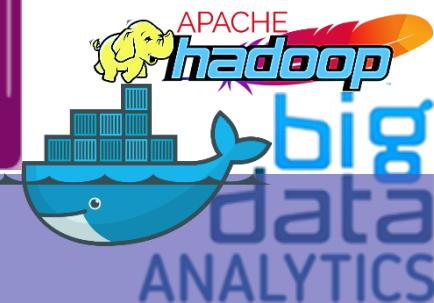
```
sohailimran@SI:~/docker-hadoop$ docker compose version
Docker Compose version v2.34.0
```

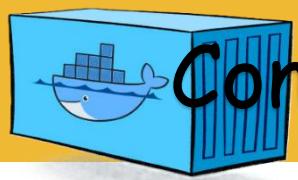
docker compose up -d

sohailimran@DESKTOP-9VAO X + ▾

- □ ×

```
sohailimran@DESKTOP-9VAOBA3:~/docker-hadoop$ docker compose up -d
WARN[0000] /home/sohailimran/docker-hadoop/docker-compose.yml: the attribute 'version' is
obsolete, it will be ignored, please remove it to avoid potential confusion
[+] Running 23/28
✓datanode Pulled                                466.9s
✓namenode Pulled                                466.6s
✓resourcemanager Pulled                           466.6s
✓nodemanager1 Pulled                            466.6s
✓historyserver Pulled                           466.6s
[+] Running 9/9
✓Network docker-hadoop_default                  Cr...      0.1s
✓Volume docker-hadoop_hadoop_historyserver    Created    0.0s
✓Volume docker-hadoop_hadoop_namenode          Created    0.0s
✓Volume docker-hadoop_hadoop_datanode          Created    0.0s
✓Container datanode                            Started   1.2s
✓Container historyserver                       Started   1.2s
✓Container nodemanager                         Started   1.3s
✓Container resourcemanager                     Starte...  1.2s
✓Container namenode                           Started   1.3s
sohailimran@DESKTOP-9VAOBA3:~/docker-hadoop$ |
```





Containers & Images

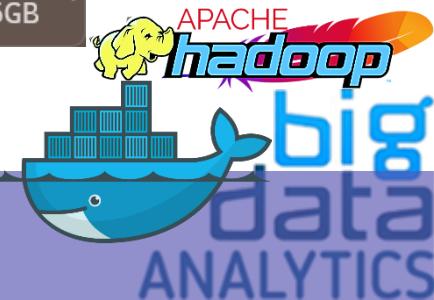


docker ps

```
sohailimran@DESKTOP-9VAO ~ docker-hadoop$ docker ps
CONTAINER ID IMAGE COMMAND CREATED STATUS
c01aabcb58dc bde2020/hadoop-historyserver:2.0.0-hadoop3.2.1-jav
utes (healthy) 8188/tcp "/entrypoint.sh /run..." 27 minutes ago Up 27 min historyserver
7a76d0a80800 bde2020/hadoop-nodemanager:2.0.0-hadoop3.2.1-jav
utes (healthy) 8042/tcp "/entrypoint.sh /run..." 27 minutes ago Up 27 min nodemanager
da398b19589c bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-jav
utes (healthy) 0.0.0.0:9000->9000/tcp, [::]:9000->9000/tcp, 0.0.0.0:9870->9870/tcp, [::]:9870->9870/tcp "/entrypoint.sh /run..." 27 minutes ago Up 27 min namenode
2aac6d26ea24 bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-jav
utes (healthy) 9864/tcp "/entrypoint.sh /run..." 27 minutes ago Up 27 min datanode
5599790f93c2 bde2020/hadoop-resourcemanager:2.0.0-hadoop3.2.1-jav
utes (healthy) 8088/tcp "/entrypoint.sh /run..." 27 minutes ago Up 27 min resourcemanager
```

docker images

```
sohailimran@DESKTOP-9VAO ~ docker-hadoop$ docker images
REPOSITORY TAG IMAGE ID CREATED SIZE
sohailimran/py-hello py2 9b05a16b76d2 7 days ago 183MB
mysql 5.7 4bc6bc963e6d 23 months ago 689MB
bde2020/hadoop-nodemanager 2.0.0-hadoop3.2.1-jav
bde2020/hadoop-resourcemanager 2.0.0-hadoop3.2.1-jav
bde2020/hadoop-namenode 2.0.0-hadoop3.2.1-jav
bde2020/hadoop-historyserver 2.0.0-hadoop3.2.1-jav
bde2020/hadoop-datanode 2.0.0-hadoop3.2.1-jav
```





APACHE
hadoop

localhost:9870



localhost:9870/dfshealth.html#tab-overview

Hadoop

Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Overview 'namenode:9000' (active)

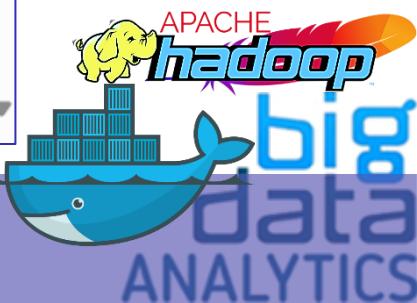
Started:	Thu Nov 27 00:10:30 +0500 2025
Version:	3.2.1, rb3cbbb467e22ea829b3808f4b7b01d07e0bf3842
Compiled:	Tue Sep 10 20:56:00 +0500 2019 by rohithsharmaks from branch-3.2.1
Cluster ID:	CID-f990dd03-1b7b-43b2-9690-44bfb37b5719
Block Pool ID:	BP-549882162-172.18.0.6-1764163124836

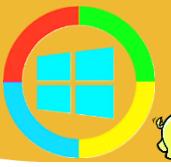
Summary

Security is off.

Safemode is off.

16 files and directories, 7 blocks (7 replicated blocks, 0 erasure coded block groups) = 23 total filesystem object(s).





APACHE
hadoop

dfs status

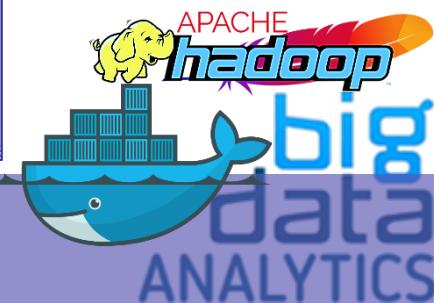


localhost:9870/dfshealth.html#tab-overview

Heap Memory used 144.61 MB of 312 MB Heap Memory. Max Heap Memory is 1.7 GB.

Non Heap Memory used 53.65 MB of 55 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.

Configured Capacity:	1006.85 GB
Configured Remote Capacity:	0 B
DFS Used:	92 KB (0%)
Non DFS Used:	4.25 GB
DFS Remaining:	951.39 GB (94.49%)
Block Pool Used:	92 KB (0%)
DataNodes usages% (Min/Median/Max/stdDev):	0.00% / 0.00% / 0.00% / 0.00%
Live Nodes	1 (Decommissioned: 0, In Maintenance: 0)
Dead Nodes	0 (Decommissioned: 0, In Maintenance: 0)
Decommissioning Nodes	0
Entering Maintenance Nodes	0
Total Datanode Volume Failures	0 (0 B)
Number of Under-Replicated Blocks	7
Number of Blocks Pending Deletion (including replicas)	0





APACHE
hadoop

Storage status



localhost:9870/dfshealth.html#tab-overview

Block Deletion Start Time	Thu Nov 27 00:10:30 +0500 2025
Last Checkpoint Time	Thu Nov 27 00:10:31 +0500 2025
Enabled Erasure Coding Policies	RS-6-3-1024k

NameNode Journal Status

Current transaction ID: 65

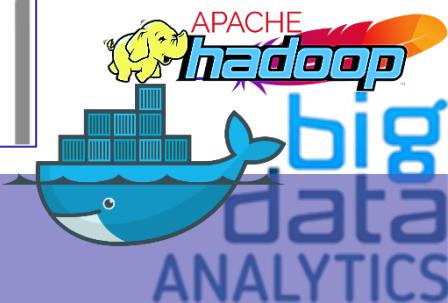
Journal Manager	State
FileJournalManager(root=/hadoop/dfs/name)	EditLogFileOutputStream(/hadoop/dfs/name/current/edits_inprogress_-1)

NameNode Storage

Storage Directory	Type	State
/hadoop/dfs/name	IMAGE_AND_EDITS	Active

DFS Storage Types

Storage Type	Configured Capacity	Capacity Used	Capacity Remaining	Block Pool Used	Nodes In Service
DISK	1006.85 GB	92 KB (0%)	951.39 GB (94.49%)	92 KB	1





namenode status



```
docker inspect -f '{{ .Config.Hostname }}' namenode
```

```
sohailimran@DESKTOP-9VAOBA3:~$ docker inspect -f '{{ .Config.Hostname }}' namenode
e89ff71ebea5
```

```
docker inspect namenode --format='{{json .NetworkSettings.Networks}}'
```

```
sohailimran@DESKTOP-9VAOBA3:~$ docker inspect namenode --format='{{json .NetworkSettings.Networks}}'
{"docker-hadoop_default":{"IPAMConfig":null,"Links":null,"Aliases":["namenode","namenode"],"MacAddress":"6a:2b:f4:60:2a:fb","DriverOpts":null,"GwPriority":0,"NetworkID":"21e7d0f80f5938110c3ffe1c8eae9c9bd5c5cdd794717c8f722271afe72dfb6e","EndpointID":"f4399829155c8264773635bbe380080839ba3f1544bd16ecb44d8ae60a6cfcb4","Gateway":"172.18.0.1","IPAddress":"172.18.0.2","IPPrefixLen":16,"IPv6Gateway":"","GlobalIPv6Address":"","GlobalIPv6PrefixLen":0,"DNSNames":["namenode","e89ff71ebea5"]}}
```

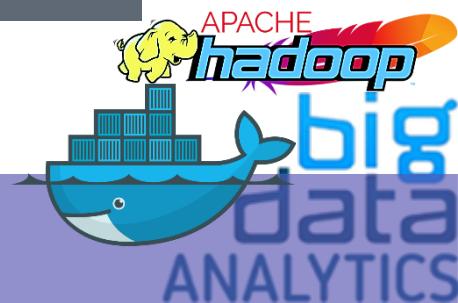
```
docker inspect -f '{{range .NetworkSettings.Networks}}{{.IPAddress}}{{end}}' namenode
```

```
sohailimran@DESKTOP-9VAOBA3:~$ docker inspect -f '{{range .NetworkSettings.Networks}}{{.IPAddress}}{{end}}' namenode
172.18.0.2
```

```
docker ps --format "table {{.Names}}\t{{.Ports}}"
```

```
sohailimran@DESKTOP-9VAOBA3:~/docker-hadoop$ docker ps --format "table {{.Names}}\t{{.Ports}}"
NAMES          PORTS
historyserver  8188/tcp
nodemanager    8042/tcp
resourcemanager 8088/tcp
namenode        0.0.0.0:9000->9000/tcp, [::]:9000->9000/tcp, 0.0.0.0:9864->9864/tcp, [::]:9864->9864/tcp,
                0.0.0.0:9870->9870/tcp, [::]:9870->9870/tcp
datanode        0.0.0.0:9865->9864/tcp, [::]:9865->9864/tcp
```

- namenode : ports 9000, 9864, 9870 exposed
- datanode : port 9865 (host) → 9864 (container) exposed



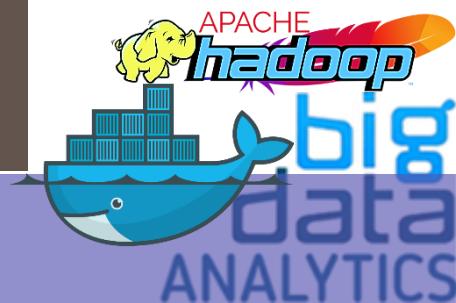


dfsadmin -report



docker exec -it namenode hdfs dfsadmin -report

```
sohailimran@DESKTOP-9VA0BA3:~/docker-hadoop$ docker exec -it namenode hdfs dfsadmin -report
Configured Capacity: 1081101176832 (1006.85 GB)
Present Capacity: 1021521588224 (951.37 GB)
DFS Remaining: 1021519667200 (951.36 GB)
DFS Used: 1921024 (1.83 MB)
DFS Used%: 0.00%
Replicated Blocks:
    Under replicated blocks: 48
    Blocks with corrupt replicas: 0
    Missing blocks: 0
    Missing blocks (with replication factor 1): 0
    Low redundancy blocks with highest priority to recover: 48
    Pending deletion blocks: 0
Erasure Coded Block Groups:
    Low redundancy block groups: 0
    Block groups with corrupt internal blocks: 0
    Missing block groups: 0
    Low redundancy blocks with highest priority to recover: 0
    Pending deletion blocks: 0
-----
Live datanodes (1):
Name: 172.18.0.6:9866 (datanode.docker-hadoop_default)
Hostname: 0dd4b4a533d5
Decommission Status : Normal
Configured Capacity: 1081101176832 (1006.85 GB)
DFS Used: 1921024 (1.83 MB)
Non DFS Used: 4587233280 (4.27 GB)
DFS Remaining: 1021519667200 (951.36 GB)
DFS Used%: 0.00%
DFS Remaining%: 94.49%
Configured Cache Capacity: 0 (0 B)
Cache Used: 0 (0 B)
Cache Remaining: 0 (0 B)
Cache Used%: 100.00%
Cache Remaining%: 0.00%
Xceivers: 1
Last contact: Sat Nov 29 05:56:25 UTC 2025
Last Block Report: Sat Nov 29 04:16:58 UTC 2025
Num of Blocks: 48
```





cp to docker



docker exec namenode ls /

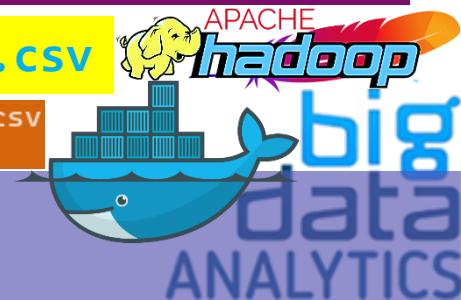
```
sohailimran@DESKTOP-9VAOBA3:~$ docker exec namenode ls /
KEYS
bin
boot
dev
entrypoint.sh
etc
hadoop
hadoop-data
home
lib
lib64
media
mnt
opt
proc
root
```

docker cp /mnt/d/DataSets/pdVsPySp.txt namenode:/home/pdVsPySp.txt

```
sohailimran@DESKTOP-9VAOBA3:~$ docker cp /mnt/d/DataSets/pdVsPySp.txt namenode:/home/pdVsPySp.txt
Successfully copied 1.54kB to namenode:/home/pdVsPySp.txt
```

docker cp /mnt/d/DataSets/salgrade.csv namenode:/home/salgrade.csv

```
sohailimran@DESKTOP-9VAOBA3:~$ docker cp /mnt/d/DataSets/salgrade.csv namenode:/home/salgrade.csv
Successfully copied 2.05kB to namenode:/home/salgrade.csv
```





docker exec -it namenode ls -lh /home/salgrade.csv

```
sohailimran@DESKTOP-9VAOBA3:~$ docker exec -it namenode ls -lh /home/salgrade.csv
-rwxrwxrwx 1 1000 1000 81 Jan 21 2024 /home/salgrade.csv
```

docker exec -it namenode hdfs dfs -mkdir /user

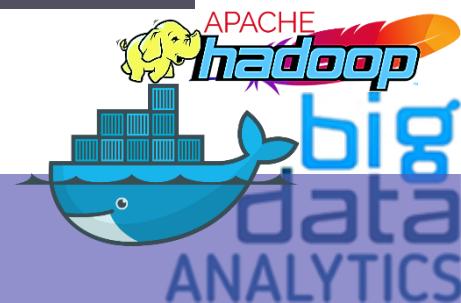
```
sohailimran@DESKTOP-9VAOBA3:~$ docker exec -it namenode hdfs dfs -mkdir /user
sohailimran@DESKTOP-9VAOBA3:~$ docker exec -it namenode hdfs dfs -ls /
Found 2 items
drwxr-xr-x  - root supergroup          0 2025-11-30 07:35 /rmstate
drwxr-xr-x  - root supergroup          0 2025-11-30 17:39 /user
```

docker exec -it namenode hdfs dfs -put /home/salgrade.csv /user/

```
sohailimran@DESKTOP-9VAOBA3:~$ docker exec -it namenode hdfs dfs -put /home/salgrade.csv /user/
2025-11-30 17:45:42,164 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted
= false, remoteHostTrusted = false
```

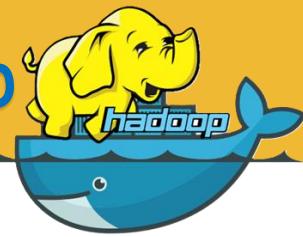
docker exec -it namenode hdfs dfs -ls /user/

```
sohailimran@DESKTOP-9VAOBA3:~$ docker exec -it namenode hdfs dfs -ls /user/
Found 1 items
-rw-r--r--  3 root supergroup          81 2025-11-30 17:45 /user/salgrade.csv
```





using Hadoop



docker context use desktop-linux



```
PS C:\Users\DELL> docker context use desktop-linux  
desktop-linux  
Current context is now "desktop-linux"
```



```
docker exec namenode env | findstr HADOOP_HOME
```

```
PS C:\Users\DELL> docker exec namenode env | findstr HADOOP_HOME  
HADOOP_HOME=/opt/hadoop-3.2.1
```



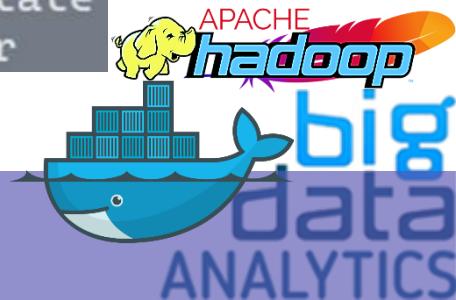
```
docker exec -it namenode bash
```

```
PS C:\Users\DELL> docker exec -it namenode bash
```



```
hdfs dfs -ls /
```

```
root@e89ff71ebea5:/# hdfs dfs -ls /  
Found 2 items  
drwxr-xr-x  - root supergroup          0 2025-11-26 13:18 /rmstate  
drwxr-xr-x  - root supergroup          0 2025-11-26 13:25 /user
```





put to hdfs



```
hdfs dfs -mkdir /input  
hdfs dfs -mkdir /output
```

```
root@e89ff71ebea5:/# hdfs dfs -mkdir /input  
root@e89ff71ebea5:/# hdfs dfs -mkdir /output
```

```
hdfs dfs -ls /
```

```
root@e89ff71ebea5:/# hdfs dfs -ls /
```

Found 4 items

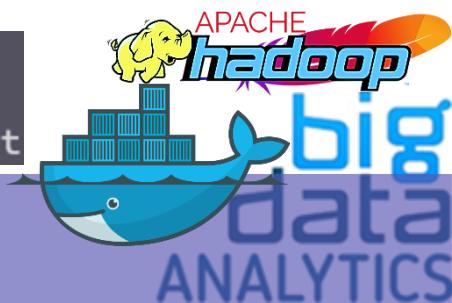
drwxr-xr-x	- root	supergroup	0 2025-11-27 13:09	/input
drwxr-xr-x	- root	supergroup	0 2025-11-27 13:09	/output
drwxr-xr-x	- root	supergroup	0 2025-11-26 13:18	/rmstate
drwxrwxrwx	- drillbits	supergroup	0 2025-11-26 20:47	/user

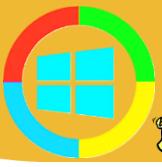
```
hdfs dfs -put /home/pdVsPySp.txt /input
```

```
root@e89ff71ebea5:/# hdfs dfs -put /home/pdVsPySp.txt /input  
2025-11-28 20:31:14,530 INFO sasl.SaslDataTransferClient: SASL encryption trust  
check: localHostTrusted = false, remoteHostTrusted = false
```

```
hdfs dfs -ls /input
```

```
root@e89ff71ebea5:/# hdfs dfs -ls /input  
Found 1 items  
-rw-r--r-- 3 root supergroup 0 2025-11-28 07:59 /input/pdVsPySp.txt
```





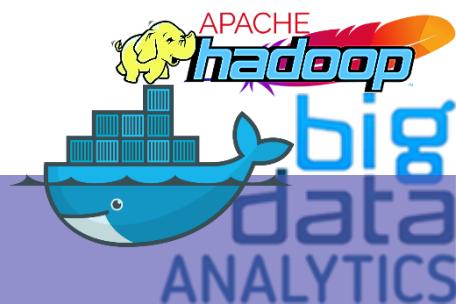
mapreduce



```
hadoop jar /opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.1.jar wordcount /input /output/wordcount
```

```
root@e89ff71ebea5:/# hadoop jar /opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.1.jar wordcount /input /output/wordcount
2025-11-28 20:37:18,644 INFO client.RMProxy: Connecting to ResourceManager at resourcemanager/172.18.0.2:8032
2025-11-28 20:37:18,971 INFO client.AHSProxy: Connecting to Application History server at historyserver/172.18.0.4:10200
2025-11-28 20:37:19,163 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/root/.staging/job_1764359819985_0002
2025-11-28 20:37:19,269 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2025-11-28 20:37:19,384 INFO input.FileInputFormat: Total input files to process : 1
2025-11-28 20:37:19,430 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2025-11-28 20:37:19,945 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2025-11-28 20:37:19,957 INFO mapreduce.JobSubmitter: number of splits:1
2025-11-28 20:37:20,240 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2025-11-28 20:37:20,254 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1764359819985_0002
2025-11-28 20:37:20,254 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-11-28 20:37:20,418 INFO conf.Configuration: resource-types.xml not found
2025-11-28 20:37:20,419 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2025-11-28 20:37:20,764 INFO impl.YarnClientImpl: Submitted application application_1764359819985_0002
2025-11-28 20:37:20,822 INFO mapreduce.Job: The url to track the job: http://resourcemanager:8088/proxy/application_1764359819985_0002/
2025-11-28 20:37:20,822 INFO mapreduce.Job: Running job: job_1764359819985_0002
2025-11-28 20:37:26,927 INFO mapreduce.Job: Job job_1764359819985_0002 running in uber mode : false
2025-11-28 20:37:26,929 INFO mapreduce.Job: map 0% reduce 0%
2025-11-28 20:37:31,991 INFO mapreduce.Job: map 100% reduce 0%
2025-11-28 20:37:36,016 INFO mapreduce.Job: map 100% reduce 100%
2025-11-28 20:37:36,025 INFO mapreduce.Job: Job job_1764359819985_0002 completed successfully
2025-11-28 20:37:36,084 INFO mapreduce.Job: Counters: 54
```

```
Peak Reduce Virtual memory (bytes)=8453066752
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=2665
File Output Format Counters
  Bytes Written=1755
```





WORD Count



hdfs dfs -ls /output

```
root@e89ff71ebea5:/# hdfs dfs -ls /output
Found 1 items
drwxr-xr-x  - root supergroup          0 2025-11-28 10:32 /output/wordcount
```

hdfs dfs -get /output/wordcount /home/wordcount

```
root@e89ff71ebea5:/# hdfs dfs -get /output/wordcount /home/wordcount
2025-11-28 20:56:09,673 INFO sasl.SaslDataTransferClient: SASL encryption
trust check: localHostTrusted = false, remoteHostTrusted = false
```

ls /home

```
sohailimran@DESKTOP-9VAOBA3:~$ ls /home
salgrade.csv  sohailimran
```



docker cp namenode:/home/wordcount /mnt/d/DataSets/wordcount

```
sohailimran@DESKTOP-9VAOBA3:~$ docker cp namenode:/home/wordcount /mnt/d/DataSets/wordcount
Successfully copied 2.56kB to /mnt/d/DataSets/wordcount
```

